



Applied Parallel Computing LLC

<http://parallel-computing.pro>

OpenACC in R language

Dmitry Mikushin Aleksei Ivakhnenko

September 6, 2016



- R is nowadays the state-of-art “Matlab” for statistical computing
- R is high-level, e.g. implements collective operations on matrices and vectors
⇒ OpenACC for R is unlikely possible in the form it exists for C/C++/Fortran



- R is nowadays the state-of-art “Matlab” for statistical computing
- R is high-level, e.g. implements collective operations on matrices and vectors
⇒ OpenACC for R is unlikely possible in the form it exists for C/C++/Fortran

So what could be done?



- R is nowadays the state-of-art “Matlab” for statistical computing
- R is high-level, e.g. implements collective operations on matrices and vectors
⇒ OpenACC for R is unlikely possible in the form it exists for C/C++/Fortran

So what could be done?

- R supports packages development in C++ with *Rcpp*
- Why not to develop a package in C++ exposing GPU through OpenACC?

Rcpp approach at glance

1 Install the *Rcpp* package:

```
$ sudo apt-get install r-cran-rcpp
```

2 Create package skeleton in R:

```
> library(Rcpp)
> Rcpp::Rcpp.package.skeleton("MarkovChain")
```

3 Add a C++ class implementation to *src/* folder:

```
$ ls src/
Makefile
Makefile.inc
MarkovChain.cpp
MarkovChainCPU.cpp
MarkovChainImpl.h
MarkovChainOpenACC.cpp
Rstreambuf.h
```

4 Build package library using R and custom Makefile:

```
$ make
R CMD INSTALL .
* installing *source* package "MarkovChain ..."
...
* DONE (MarkovChain)
```

5 Use the created package in R:

```
$ R

> library(MarkovChain)
> m = new("MarkovChain", c("a", "a", "b"), TRUE)
Using OpenACC implementation of MarkovChain
> m$GetTransitionMatrix()
      [,1] [,2]
[1,]    1    0
[2,]    0    0
```

Example port of package code into OpenACC kernel

```
NumericMatrix p(Rank, Rank);  
...  
NumericVector s(Rank, 0.);  
...  
for (int i = 0; i < Rank; i++)  
{  
  if (s[i] == 0)  
    continue;  
  for (int j = 0; j < Rank; j++)  
    p(i,j) = p(i,j) / s[i];  
}  
  
this->TransitionMatrix = p;
```



```
...  
double* parr = REAL(p);  
double* sarr = REAL(s);  
...  
openacc_kernel(sarr, parr, Rank);  
...  
static void openacc_kernel(double* sarr, double* parr, int Rank)  
{  
  #pragma acc kernels loop independent copyin(sarr[0:Rank]) copy(parr[0:Rank*Rank])  
  for (int i = 0; i < Rank; i++)  
  {  
    #pragma acc loop independent  
    for (int j = 0; j < Rank; j++)  
    {  
      if (sarr[i] != 0)  
        parr[i + j * Rank] /= sarr[i];  
    }  
  }  
}
```

- PGI OpenACC is not perfect for C++: suffers from crashes when adding directives into complex C++ class methods

Solution: Extract OpenACC-ported loops into separate static functions:

```
...  
openacc_kernel(sarr, parr, Rank);  
...  
static void openacc_kernel(double* sarr, double* parr, int Rank)  
{  
    ...  
}
```

- C++ data containers need a lot of work to get ported to GPUs

Solution: Much easier to switch to plain pointers where possible:

```
...  
double* parr = REAL(p);  
double* sarr = REAL(s);  
...
```

Check whether OpenACC is actually running

Add `PGI_ACC_TIME=1` to have the profiling info after exiting from R console:

```
$ PGI_ACC_TIME=1 R

R version 3.0.2 (2013-09-25) — "Frisbee Sailing"

> library(MarkovChain)
> m = new("MarkovChain", c("a", "a", "b"), TRUE)
Using OpenACC implementation of MarkovChain
> q()
Save workspace image? [y/n/c]: n

Accelerator Kernel Timing data
/usr/lib/R/site-library/Rcpp/include/Rcpp/routines.h
_ZN44_INTERNAL_22_MarkovChainOpenACC_cpp_c334669514openacc_kernelEPdS0_i NVIDIA devicenum=0
time(us): 37
9: compute region reached 1 time
  14: kernel launched 1 time
    grid: [1x2] block: [128]
    elapsed time(us): total=41 max=41 min=41 avg=41
9: data region reached 2 times
  9: data copyin transfers: 2
    device time(us): total=26 max=16 min=10 avg=13
  19: data copyout transfers: 1
    device time(us): total=11 max=11 min=11 avg=11
```